

# Fundamentals in R: Extended Functionality

August 12, 2014







# Module 5 - Extended Functionality

In this module we'll introduce how to extend the basic functionality of R by installing and loading additional packages. We'll introduce the most popular data visualization package in R.

The objectives are

- install and load R packages
- get a feel for ggplot2
- put your acquired skills to good use





# Extending R

All of the functionality we've explored thus far ships standard with R. The beauty of R lies in it being *open source* and *extendable*. Additional functionality is provided by functions shipped in packages. Packages must be installed once (for every R version install) and loaded (every time you start R).

Identifying packages for download:

[Search rseek.org](https://rseek.org) [Browse CRAN Task Views](#)



# Exercise

- Identify a business/analytics problem you would like to tackle with R  
(examples: importing data into R from Excel, running a logistic regression in R, R with Google Motion Charts)
- Search the community
  - Who else has been working on this? [r-bloggers.com](http://r-bloggers.com)
  - What package(s) could prove useful? [Search rseek.org](http://Search.rseek.org)
  - Can you leverage code that's already been written?
- Identify a package to download



# Load Your Work

```
load("data/Module3_songData.RData")  
load("data/Module4_songSubset3.RData")
```





# Installing, Loading, Help

```
install.packages("ggplot2")  # Install a package
```

```
library("ggplot2")  # Load a package
```

```
load("data/Module4_songSubset3.RData")
```

```
help(package = "ggplot2")  # Get help on a package
```

```
help("map_data")  # Get help for a specific function
```

```
# ... or ?map_data
```

# ggplot2



ggplot2 is the most popular package for data visualization in R. It's syntax may be very different from anything you've worked with in the past but it is precisely what gives ggplot it flexibility and overall ease of use.

To see examples of how others have used this package/function, a recommended resource is [www.rseek.org](http://www.rseek.org) search. Here's an [example](#).







# ggplot2

```
mapWorld <- map_data(map = "world")
```

```
# What does map_data() function do?
```

```
# ?map_data
```



# ggplot2

*# What type of object is mapWorld?*

```
str(mapWorld)
```

```
## 'data.frame':    25553 obs. of  6 variables:
## $ long      : num  -133 -132 -132 -132 -130 ...
## $ lat       : num   58.4 57.2 57 56.7 56.1 ...
## $ group     : num   1 1 1 1 1 1 1 1 1 1 ...
## $ order     : int   1 2 3 4 5 6 7 8 9 10 ...
## $ region    : chr   "Canada" "Canada" "Canada" "Canada" ...
## $ subregion: chr   NA NA NA NA ...
```

*# What type of object is mapWorld?*

```
head(mapWorld)
```

```
##      longitude lat group order region subregion
## 1 -133.4 58.42      1      1 Canada      <NA>
```



# ggplot2


ggplot2 can take the data frame generated by `map_data` and plot it as polygons using latitude and longitude coordinates as well as group information.

```
songMap <- ggplot() + geom_polygon(data = mapWorld, aes(x = long, y = lat,  
  fill = "white"))
```

```
songMap
```



# ggplot2



# ggplot2




We can add data to our map. Each point will represent an observation in our song data frame (one song). The size of the point will be relative to the hfRatio.

```
songMap + geom_point(data = songSubset3, aes(x = artist.longitude, y = a  
size = hfRatio), alpha = I(0.5))
```





# ggplot2



How can we improve this map?



# Putting it all together

Putting it all together, let's plot average hfRatio in the US by state – identifying states with up and coming artists.





# Exercise

From `songSubset3`, remove any rows with NA's or NaN's for variables we'll be plotting. You will need to:

- 1 Identify which variables we need to screen for NA's and NaN's
- 2 Use `!is.na()` to identify observations in our data (`songSubset3`) where those variables *are not NA or NaN*
- 3 Use `[]` to extract a subset data frame (call it `songState`) based on the information obtained in 2.







# Solution

```
songState <- songSubset3[!is.na(songSubset3$artist.hotttnesss) & !is.na  
  !is.na(songSubset3$song.hotttnesss) & !is.na(songSubset3$artist.lat  
  !is.na(songSubset3$artist.longitude), ]
```

```
dim(songSubset3)
```

```
## [1] 3518 14
```

```
dim(songState)
```

```
## [1] 2210 14
```



# Solution

*# Note, we can identify (remove) all rows where observations of ANY of the  
# variables are NA using complete.cases (na.omit).*

```
dim(na.omit(songSubset3))
```

```
## [1] 1435    14
```

```
dim(songSubset3[complete.cases(songSubset3), ])
```

```
## [1] 1435    14
```





# Putting it all together

We want to plot hotness and familiarity variables by state. Which variables can we use to do that?

```
View(songState)
```





# Putting it all together

What are the next steps for plotting average hotness and familiarity variables by US state?





# Putting it all together

- 1 Subset data to only include US
- 2 Aggregate hotness and familiarity variables by State



# Exercise



- 1 Subset `songState` to only include rows where `geoCountry` is "United States". Save this as an object called `songStateSub`. Hint: use `[]` or `subset()`. Under which state do most of our observations fall? How many observations?





# Solution

```
songStateSub <- subset(songState, geoCountry == "United States")
```

```
summary(songStateSub)
```

```
## artist.latitude artist.longitude      geoCountry      geoState
## Min.   :19.6    Min.   : -155.4   United States:1448   California:263
## 1st Qu.:34.0    1st Qu.: -100.1   Afghanistan : 0    New York :217
## Median :37.3    Median : -87.7    Argentina   : 0    Texas   :120
## Mean   :37.1    Mean   : -92.0    Australia   : 0    Tennessee: 68
## 3rd Qu.:40.7    3rd Qu.: -77.3    Austria     : 0    Illinois : 53
## ...
```

# Exercise



- 2 Aggregate songState hfRatio variable by calculation the mean over the geoState variable. Call this data set songStateAgg.





# Solution



```
songStateAgg <- aggregate(hfRatio ~ geoState, data = songStateSub, mean.)
```



# Putting it all together

To generate the visualization of, say, average hfRatio by state, we will walk through a similar process as before.

As before, we use `map_data` to obtain a data frame which `ggplot2` can use to plot. This time, we obtain the state map which corresponds to US states.

```
usStateMap <- map_data("state")  
head(usStateMap)
```

```
##      long   lat group order  region subregion  
## 1 -87.46 30.39     1     1 alabama      <NA>  
## 2 -87.48 30.37     1     2 alabama      <NA>  
## 3 -87.53 30.37     1     3 alabama      <NA>  
## 4 -87.53 30.33     1     4 alabama      <NA>  
## 5 -87.57 30.33     1     5 alabama      <NA>  
## 6 -87.59 30.33     1     6 alabama      <NA>
```



# Exercise

This time, however, because we are plotting the songhottness variable by state and not simply by lat/long coordinates, we need to join our data with this map data frame on state name.

Join `usStateMap` and `songStateAgg` on the state names.





# Solution

```
mapData <- merge(songStateAgg, usStateMap, by.x = "geoState", by.y = "r
```

This merge failed. Why? What do we need to ensure a successful merge?

```
head(mapData)
```

```
## [1] geoState hfRatio long    lat      group    order    subregion  
## <0 rows> (or 0-length row.names)
```



## Solution part 2

```
head(usStateMap, 2)
```

```
##      long   lat group order  region subregion
## 1 -87.46 30.39     1     1  alabama      <NA>
## 2 -87.48 30.37     1     2  alabama      <NA>
```

```
head(songStateAgg, 2)
```

```
##   geoState hfRatio
## 1  Alabama  0.6883
## 2  Alaska   0.8668
```

```
songStateAgg$geoState <- tolower(songStateAgg$geoState)
```





## Solution part 2

```
mapData <- merge(songStateAgg, usStateMap, by.x = "geoState", by.y = "r  
  all.y = TRUE)  
head(mapData)
```

##	geoState	hfRatio	long	lat	group	order	subregion
## 1	alabama	0.6883	-87.46	30.39	1	1	<NA>
## 2	alabama	0.6883	-87.48	30.37	1	2	<NA>
## 3	alabama	0.6883	-85.10	32.64	1	119	<NA>
## 4	alabama	0.6883	-85.07	32.61	1	120	<NA>
## 5	alabama	0.6883	-85.05	32.56	1	121	<NA>
## 6	alabama	0.6883	-87.53	30.37	1	3	<NA>





# Putting it all together

Notice that in merging, our data frame got a little bit jumbled up. We'll re-order the observations based on the order variable.

```
head(usStateMap, 3)
```

```
##      long   lat group order  region subregion
## 1 -87.46 30.39     1     1 alabama      <NA>
## 2 -87.48 30.37     1     2 alabama      <NA>
## 3 -87.53 30.37     1     3 alabama      <NA>
```

```
head(mapData, 3)
```

```
##   geoState hfRatio   long   lat group order subregion
## 1  alabama  0.6883 -87.46 30.39     1     1      <NA>
## 2  alabama  0.6883 -87.48 30.37     1     2      <NA>
## 3  alabama  0.6883 -85.10 32.64     1    119      <NA>
```

```
mapData <- mapData[order(mapData$order), ]
```



# Putting it all together

Our ggplot code will be very much the same as it was before. Instead of `geom_point`, we will use `geom_polygon` as we are plotting states, not coordinates.

```
(p1 <- ggplot(data = mapData, aes(x = long, y = lat, group = group)) +  
  5))))
```


```
(p1 <- p1 + geom_path(colour = "orange", linestyle = 2))
```

```
(p1 <- p1 + scale_fill_brewer("Average Artists Hotness to Familiarity R  
  palette = "PuRd"))
```





# Putting it all together





# Putting it all together

We add state names (abbrev.) using a built in dataset containing lat/long state centers. We clean up our map, removing axes and background. We save the image to a jpeg file using ggsave().

```
states <- data.frame(state.center, state.abb)


(p1 <- p1 + geom_text(data = states, aes(x = x, y = y, label = state.abb,
size = 2))

(p1 <- p1 + theme(line = element_blank(), axis.text = element_blank(),
panel.background = element_blank()))

ggsave(filename = "p1.jpeg", width = 11, height = 8.5)
```



# Putting it all together





# Where to next?

Now that we know which states, on average, are associated with the most up & coming artists by average rating, what else might we want to investigate?





# Where to next?

How does number of observations vary across states? How can we incorporate this information into our plot?

How do the hotttness and familiarity variables vary by release year? How does this affect the ratio we calculated?





...

Thank You!



# Thank you

Revolution Analytics is the leading commercial provider of software and support for the popular open source R statistics language.

[www.revolutionanalytics.com](http://www.revolutionanalytics.com), 1.855.GET.REVO, Twitter: @RevolutionR

